

# EffPPL

A Probabilistic Programming Library using  
Effect Handlers in Multicore OCaml

**Arnhav Datar**

Guided by: KC Sivaramakrishnan and Tom Kelly



# Agenda

Objective

Effect Handlers

Probabilistic Programming

EffPPL, Inference Algorithm

Algorithmic Differentiation

Application and Results

Evaluation

Q&A

25min

5 min

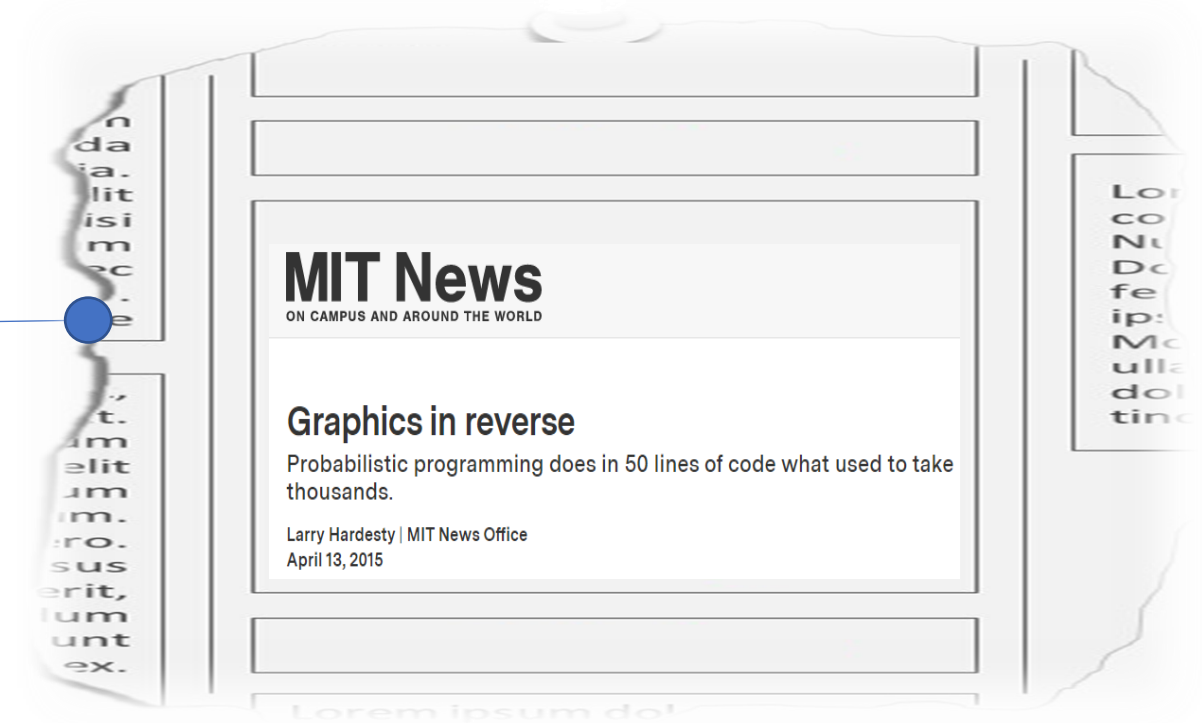


# Objective - EffPPL

With the recent rise of Probabilistic Programming Languages (PPLs) in various domains, we proposed to build a library as follows.

- Concurrent
- Probabilistic Programming Library
- in Multicore OCaml
- made using Effect Handlers.

Deep probabilistic programming (DPP) combines three fields: Bayesian statistics and machine learning, deep learning, and probabilistic programming. Many applications in trading and risk management involve uncertainty quantification such as portfolio optimization and risk estimation. DPP provides a general purpose Bayesian framework for fitting high dimensional models to large datasets to provide a richer set of statistical properties and more explanatory power.



Another benefit is that PPLs reduce the amount of manually written code for a particular inference problem, facilitating the task and minimizing the risk of inadvertently introducing errors, biases or inaccuracies. Our verification experiments suggest that the light-weight PPL implementations of ClaDS1 and ClaDS2 provide more accurate computation of likelihoods than the thousands of lines of code developed originally for these models.

# Objective

With the recent rise of Probabilistic Programming Languages(PPLs) in various domains, we present EffPPL. EffPPL is a

- Shallowly-Embedded
- Probabilistic Programming Library
- in Multicore OCaml
- made using Effect Handlers.



# Effect Handlers

- Effect handlers are a mechanism for modular programming with user-defined effects.
- Similar to defining new exception values, the user can define their own effects. These are handled with an effect handler.
- But unlike exceptions, effect handlers permit the control to go back to where the effect was performed, but also save the context necessary for going back in a data structure.
- Multicore OCaml incorporates effect handlers as a way of supporting concurrency primitives.



## Effect handler : Example

```
effect E : string

let comp () =
  print_string "0 ";
  print_string (perform E);
  print_string "3 "

let main () =
  try
    comp ()
  with effect E k ->
    print_string "1 ";
    continue k "2 ";
    print_string "4 "
```



## Effect handler : Example

effect declaration

```
effect E : string

let comp () =
  print_string "0 ";
  print_string (perform E);
  print_string "3 "

let main () =
  try
    comp ()
  with effect E k ->
    print_string "1 ";
    continue k "2 ";
    print_string "4 "
```



# Effect handler : Example

effect declaration

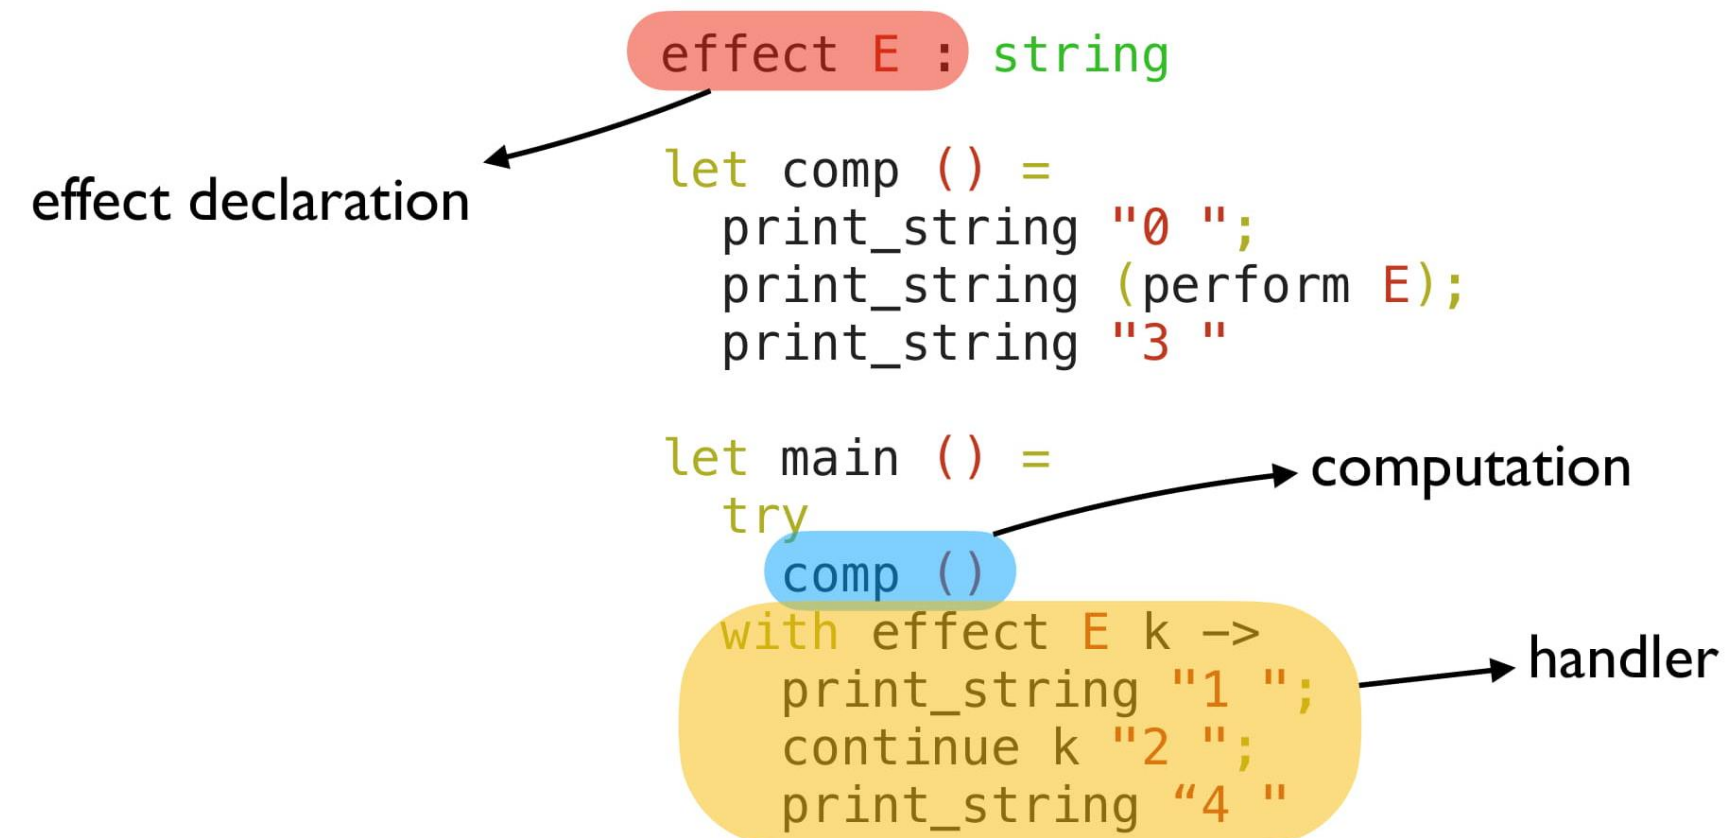
```
effect E : string
```

computation

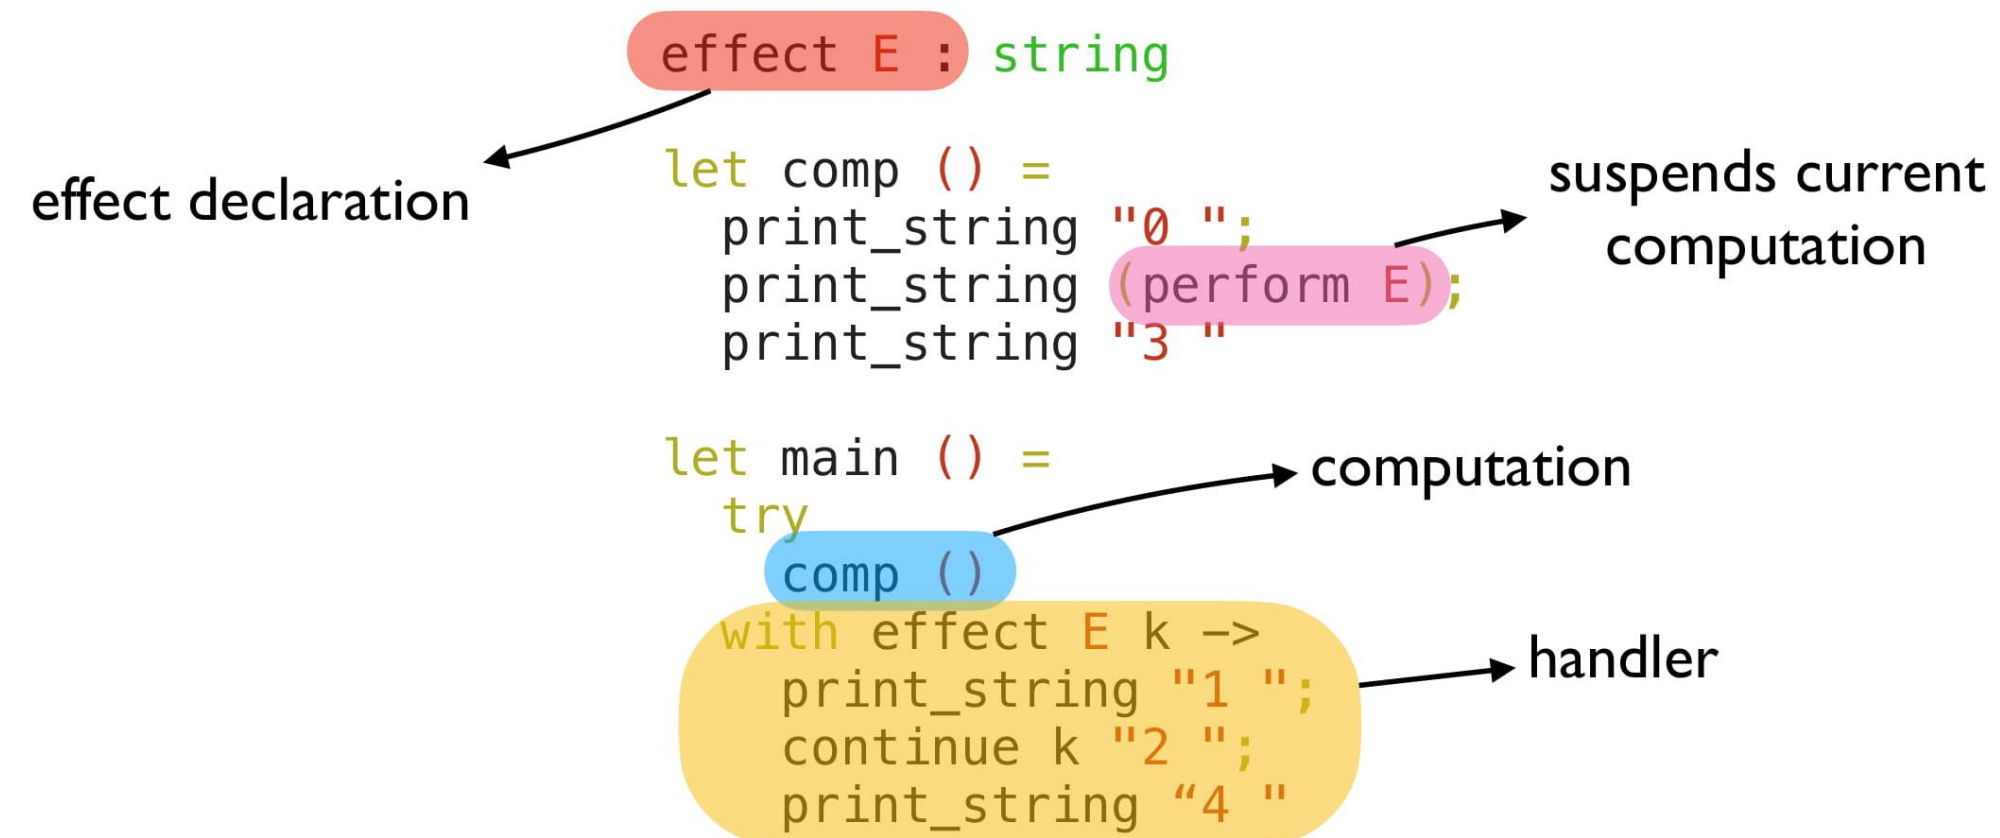
```
let comp () =  
  print_string "0 ";  
  print_string (perform E);  
  print_string "3 "  
  
let main () =  
  try  
    comp ()  
  with effect E k ->  
    print_string "1 ";  
    continue k "2 ";  
    print_string "4 "
```



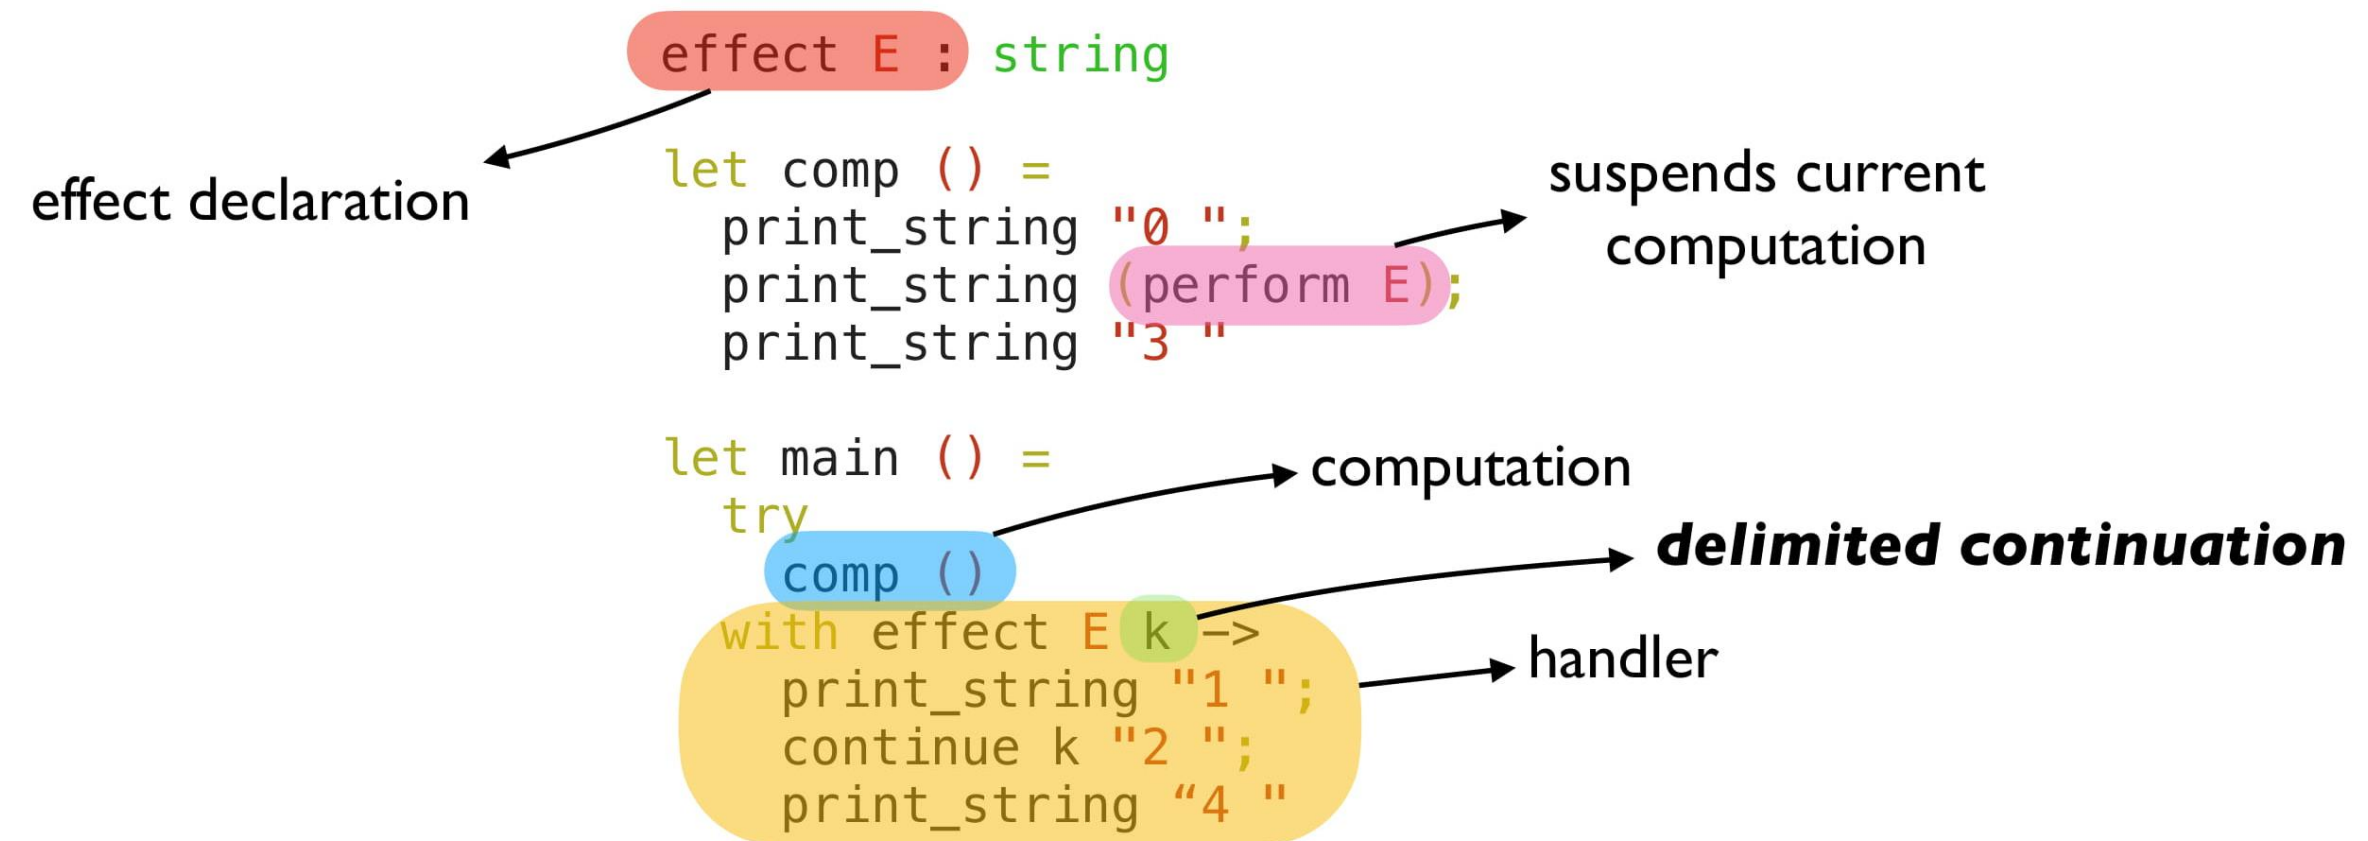
# Effect handler : Example



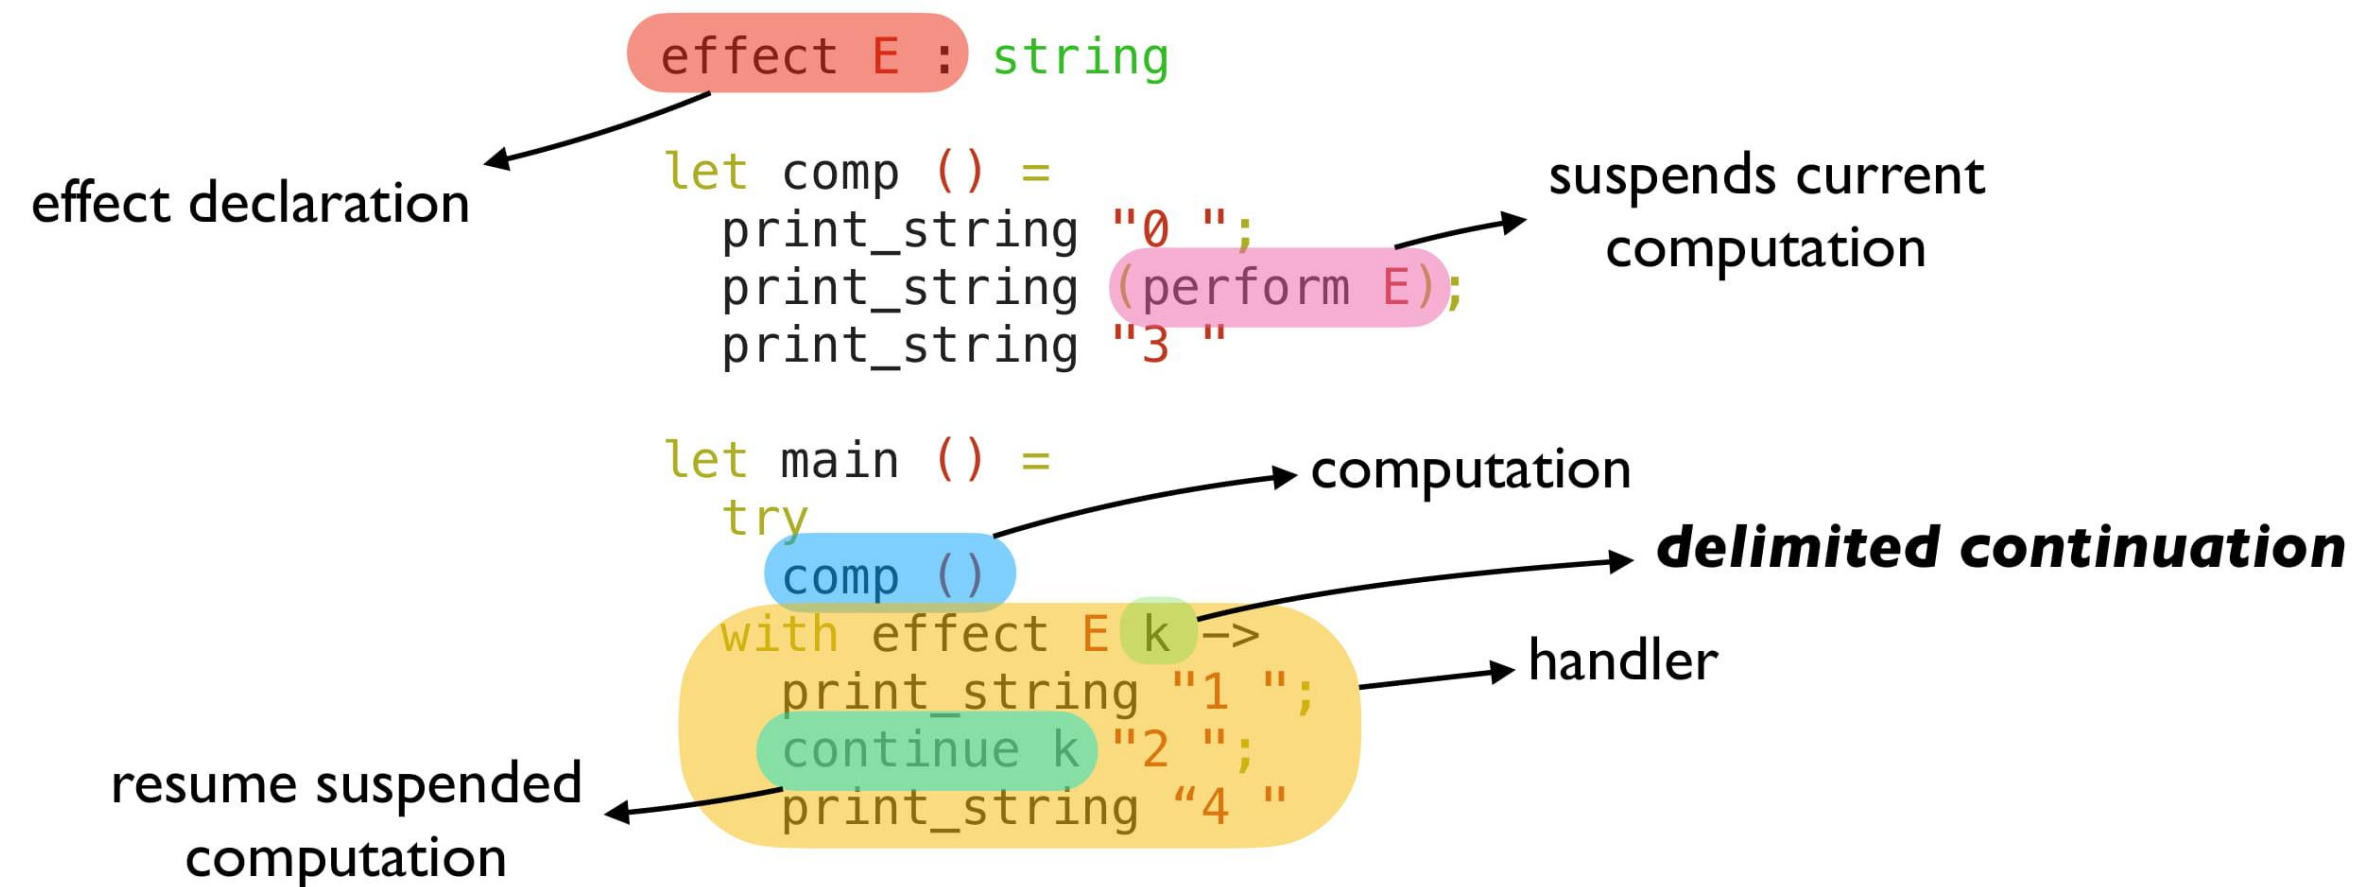
# Effect handler : Example



# Effect handler : Example



# Effect handler : Example

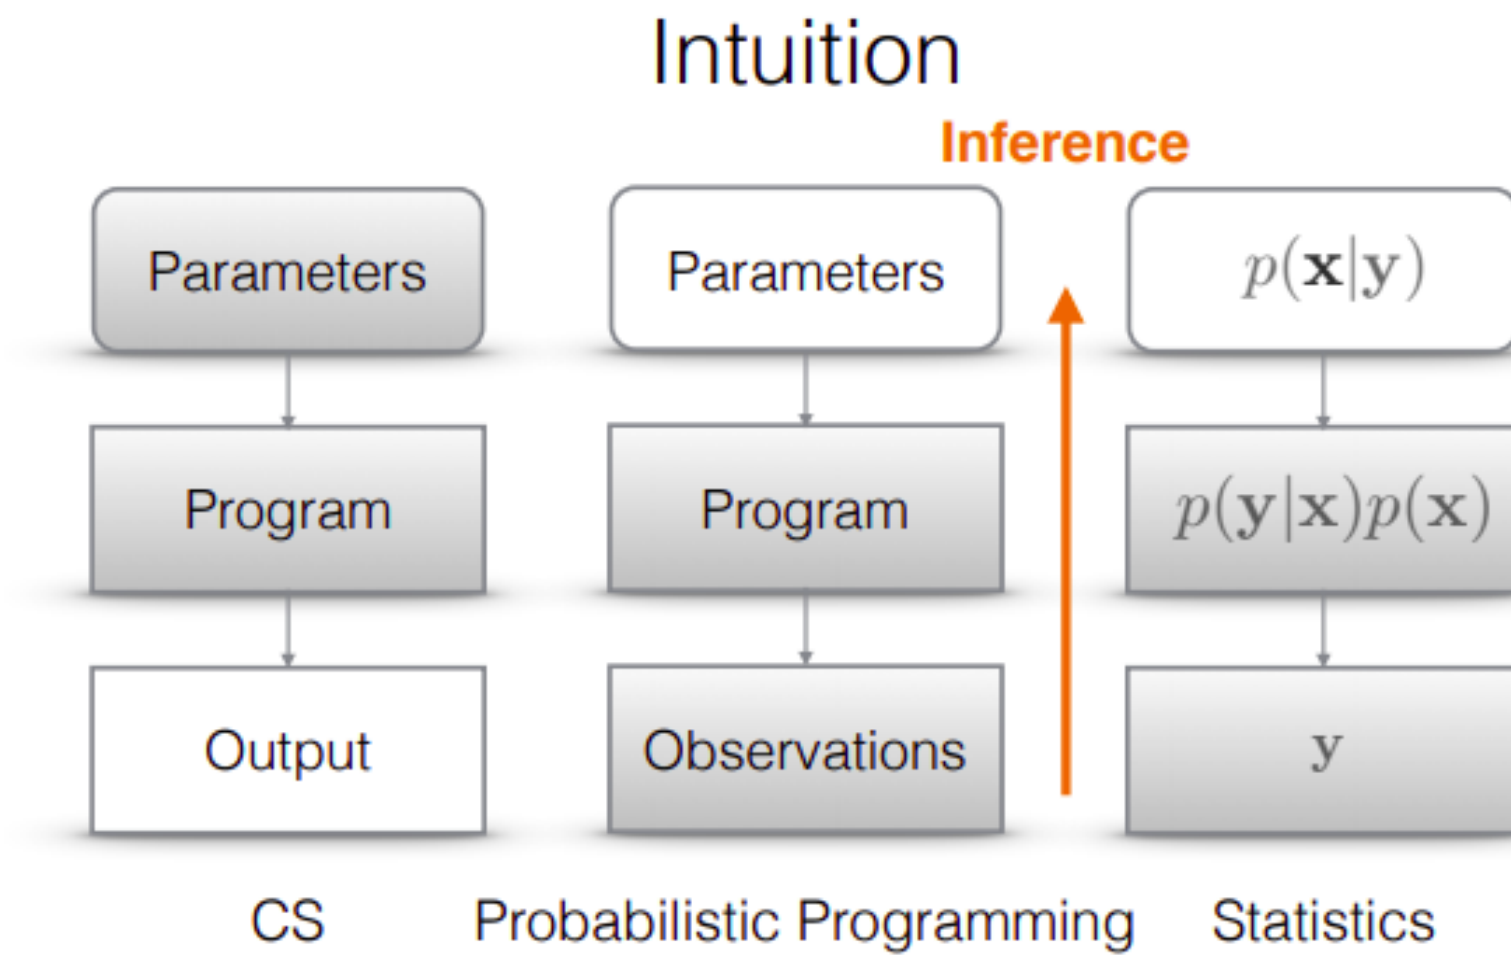


# Probabilistic Programming

- Probabilistic Programming involves the construction of inference problems and the development of corresponding evaluators, that computationally characterize the denoted conditional distribution.
- Probabilistic programming languages(PPLs) have been used for performing approximate Bayesian inference in complex generative models. PPLs allow us to create these models using simpler models and probability distributions.




# Probabilistic Programming



## A WebPPL Example


```
var binomial = function() {  
  var a = sample(Bernoulli({ p: 0.5 })))  
  var b = sample(Bernoulli({ p: 0.5 })))  
  var c = sample(Bernoulli({ p: 0.5 })))  
  return a + b + c  
}  
  
var binomialDist = Infer({ model: binomial })  
  
viz(binomialDist)
```






## A WebPPL Example

```
var binomial = function() {  
  var a = sample(Bernoulli({ p: 0.5 })))  
  var b = sample(Bernoulli({ p: 0.5 })))  
  var c = sample(Bernoulli({ p: 0.5 })))  
  return a + b + c  
}  
  
var binomialDist = Infer({ model: binomial })  
  
viz(binomialDist)
```



## A WebPPL Example

```
var funnyBinomial = function(){  
  var a = sample(Bernoulli({ p: 0.5 }))  
  var b = sample(Bernoulli({ p: 0.5 }))  
  var c = sample(Bernoulli({ p: 0.5 }))  
  factor( (a || b) ? 0 : -2)  
  return a + b + c  
}  
  
var funnyBinomialDist = Infer({ model: funnyBinomial })  
  
viz(funnyBinomialDist)
```



# Language Design

```
let sum_of_floats () =  
  let x1 = 2. in  
  let x2 = 3. in  
  x1 +. x2
```

Sum of two floats in OCaml

EffPPL



```
let sum_of_normals () =  
  let* x1 = normal 0. 1. in  
  let* x2 = normal 0. 1. in  
  x1 +. x2
```

Sum of two standard normals

For the user to have an intuitive and easily understandable interface to create the models the language tries to mirror OCaml with the only difference being using `let*` as opposed to `let`.

The language is also shallowly embedded so all other OCaml functions and primitives such as `for`, `if` etc. work within the EffPPL functions.



# Hamiltonian Monte-Carlo

---

**Algorithm 1:** Hamiltonian Monte Carlo
 

---

Given  $\theta^0, L, \epsilon, \mathcal{L}, M$ :

**for**  $m = 1$  *to*  $M$  **do**

  Sample  $r^0 \sim \mathcal{N}(0, I)$

$\theta^m \leftarrow \theta^{m-1}$

$\theta' \leftarrow \theta^{m-1}$

$r' \leftarrow r^0$

**for**  $i = 1$  *to*  $L$  **do**

$(\theta', r') \leftarrow \text{Leapfrog}(\theta', r', \epsilon)$

**end**

  With probability  $\min\{1, \frac{\exp(\mathcal{L}(\theta') - 0.5 \cdot r' \cdot r')}{\exp(\mathcal{L}(\theta^{m-1}) - 0.5 \cdot r^0 \cdot r^0)}\}$  set  $\theta^m \leftarrow \theta'$  and  $r^m \leftarrow -r'$

**end**

**function**  $\text{Leapfrog}(\theta', r', \epsilon)$  :

$r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta)$

$\theta' \leftarrow \theta' + (\epsilon) r'$

$r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta')$

**return**  $(\theta', r')$

---

# Hamiltonian Monte-Carlo

---

**Algorithm 1:** Hamiltonian Monte Carlo
 

---

 Given  $\theta^0, L, \epsilon, \mathcal{L}, M$ :

 Initial sample, iterations, step size, model, epochs
 

---

**for**  $m = 1$  *to*  $M$  **do**

 Sample  $r^0 \sim \mathcal{N}(0, I)$ 
 $\theta^m \leftarrow \theta^{m-1}$ 
 $\theta' \leftarrow \theta^{m-1}$ 
 $r' \leftarrow r^0$ 
**for**  $i = 1$  *to*  $L$  **do**

 |  $(\theta', r') \leftarrow \text{Leapfrog}(\theta', r', \epsilon)$ 
**end**

 With probability  $\min\{1, \frac{\exp(\mathcal{L}(\theta') - 0.5 \cdot r' \cdot r')}{\exp(\mathcal{L}(\theta^{m-1}) - 0.5 \cdot r^0 \cdot r^0)}\}$  set  $\theta^m \leftarrow \theta'$  and  $r^m \leftarrow -r'$ 
**end**
**function**  $\text{Leapfrog}(\theta', r', \epsilon)$  :

 $r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta)$ 
 $\theta' \leftarrow \theta' + (\epsilon) r'$ 
 $r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta')$ 

 return  $(\theta', r')$ 


---



# Hamiltonian Monte-Carlo

---

## Algorithm 1: Hamiltonian Monte Carlo

---

Given  $\theta^0, L, \epsilon, \mathcal{L}, M$ :

Initial sample, iterations, step size, model, epochs

for  $m = 1$  to  $M$  do

Sample  $r^0 \sim \mathcal{N}(0, I)$

$\theta^m \leftarrow \theta^{m-1}$

$\theta' \leftarrow \theta^{m-1}$

$r' \leftarrow r^0$

Initialising candidate  
momentum and position

for  $i = 1$  to  $L$  do

|  $(\theta', r') \leftarrow \text{Leapfrog}(\theta', r', \epsilon)$

end

With probability  $\min\{1, \frac{\exp(\mathcal{L}(\theta') - 0.5 \cdot r' \cdot r')}{\exp(\mathcal{L}(\theta^{m-1}) - 0.5 \cdot r^0 \cdot r^0)}\}$  set  $\theta^m \leftarrow \theta'$  and  $r^m \leftarrow -r'$

end

**function**  $\text{Leapfrog}(\theta', r', \epsilon)$  :

$r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta)$

$\theta' \leftarrow \theta' + (\epsilon) r'$

$r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta')$

return  $(\theta', r')$

---

# Hamiltonian Monte-Carlo

---

## Algorithm 1: Hamiltonian Monte Carlo

---

Given  $\theta^0, L, \epsilon, \mathcal{L}, M$ : Initial sample, iterations, step size, model, epochs

**for**  $m = 1$  *to*  $M$  **do**

Sample  $r^0 \sim \mathcal{N}(0, I)$

$\theta^m \leftarrow \theta^{m-1}$

$\theta' \leftarrow \theta^{m-1}$

$r' \leftarrow r^0$

Initialising candidate  
momentum and position

**for**  $i = 1$  *to*  $L$  **do**

|  $(\theta', r') \leftarrow \text{Leapfrog}(\theta', r', \epsilon)$

**end**

With probability  $\min\{1, \frac{\exp(\mathcal{L}(\theta') - 0.5 \cdot r' \cdot r')}{\exp(\mathcal{L}(\theta^{m-1}) - 0.5 \cdot r^0 \cdot r^0)}\}$  set  $\theta^m \leftarrow \theta'$  and  $r^m \leftarrow -r'$

**end**

**function**  $\text{Leapfrog}(\theta', r', \epsilon)$  :

$r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta)$

$\theta' \leftarrow \theta' + (\epsilon) r'$

$r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta')$

return  $(\theta', r')$

Numerical  
Integrator



# Hamiltonian Monte-Carlo

---

## Algorithm 1: Hamiltonian Monte Carlo

---

Given  $\theta^0, L, \epsilon, \mathcal{L}, M$ :

Initial sample, iterations, step size, model, epochs

for  $m = 1$  to  $M$  do

Sample  $r^0 \sim \mathcal{N}(0, I)$

$\theta^m \leftarrow \theta^{m-1}$

$\theta' \leftarrow \theta^{m-1}$

$r' \leftarrow r^0$

Initialising candidate  
momentum and position

for  $i = 1$  to  $L$  do

|  $(\theta', r') \leftarrow \text{Leapfrog}(\theta', r', \epsilon)$

end

With probability  $\min\{1, \frac{\exp(\mathcal{L}(\theta') - 0.5 \cdot r' \cdot r')}{\exp(\mathcal{L}(\theta^{m-1}) - 0.5 \cdot r^0 \cdot r^0)}\}$  set  $\theta^m \leftarrow \theta'$  and  $r^m \leftarrow -r'$

Metropolis  
Acceptance

end

function  $\text{Leapfrog}(\theta', r', \epsilon)$ :

$r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta)$


$\theta' \leftarrow \theta' + (\epsilon) r'$

$r' \leftarrow r' + (\epsilon/2) \nabla_{\theta} \mathcal{L}(\theta')$

return  $(\theta', r')$


Numerical  
Integrator

# Hamiltonian Monte-Carlo Example



# Hamiltonian Monte-Carlo Example

## Derivative of LogPDF of standard normal distribution




# Algorithmic Differentiation

```

| r ->
  r.d <- 1.0;
  ls := modif_der !ls r.v r.d;
  (r)
| effect (Add(a,b)) k ->
  let t = {v = a.v +. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. t.d;
  b.d <- b.d +. t.d;
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)

| effect (Mult(a,b)) k ->
  let t = {v = a.v *. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. (b.v *. t.d);
  b.d <- b.d +. (a.v *. t.d);
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)

```



$$z = x + xy$$

$$x = 3$$


$$y = 4$$

Task: To find partial derivatives of  $z$  with respect to  $x$  and  $y$

# Algorithmic Differentiation

```
| r ->
  r.d <- 1.0;
  ls := modif_der !ls r.v r.d;
  (r)
| effect (Add(a,b)) k ->
  let t = {v = a.v +. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. t.d;
  b.d <- b.d +. t.d;
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)
```

```
| effect (Mult(a,b)) k ->
  let t = {v = a.v *. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. (b.v *. t.d);
  b.d <- b.d +. (a.v *. t.d);
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)
```




# Algorithmic Differentiation

```

| r ->
  r.d <- 1.0;
  ls := modif_der !ls r.v r.d;
  (r)
| effect (Add(a,b)) k ->
  let t = {v = a.v +. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. t.d;
  b.d <- b.d +. t.d;
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)

| effect (Mult(a,b)) k ->
  let t = {v = a.v *. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. (b.v *. t.d);
  b.d <- b.d +. (a.v *. t.d);
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)

```



# Algorithmic Differentiation

```
| r ->
```


```
r.d <- 1.0;  
ls := modif_der !ls r.v r.d;  
(r)
```

```
| effect (Add(a,b)) k ->
```

```
let t = {v = a.v +. b.v; d = 0.; m=1.} in  
ignore (continue k t);  
a.d <- a.d +. t.d;  
b.d <- b.d +. t.d;  
ls := modif_der !ls a.v a.d;  
ls := modif_der !ls b.v b.d;  
(x)
```

```
| effect (Mult(a,b)) k ->
```

```
let t = {v = a.v *. b.v; d = 0.; m=1.} in  
ignore (continue k t);  
a.d <- a.d +. (b.v *. t.d);  
b.d <- b.d +. (a.v *. t.d);  
ls := modif_der !ls a.v a.d;  
ls := modif_der !ls b.v b.d;  
(x)
```






# Algorithmic Differentiation

```

| r ->
  r.d <- 1.0;
  ls := modif_der !ls r.v r.d;
  (r)
| effect (Add(a,b)) k ->
  let t = {v = a.v +. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. t.d;
  b.d <- b.d +. t.d;
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)

| effect (Mult(a,b)) k ->
  let t = {v = a.v *. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. (b.v *. t.d);
  b.d <- b.d +. (a.v *. t.d);
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)

```



# Algorithmic Differentiation


```

| r ->
  r.d <- 1.0;
  ls := modif_der !ls r.v r.d;
  (r)

| effect (Add(a,b)) k ->
  let t = {v = a.v +. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. t.d;
  b.d <- b.d +. t.d;
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)


| effect (Mult(a,b)) k ->
  let t = {v = a.v *. b.v; d = 0.; m=1.} in
  ignore (continue k t);
  a.d <- a.d +. (b.v *. t.d);
  b.d <- b.d +. (a.v *. t.d);
  ls := modif_der !ls a.v a.d;
  ls := modif_der !ls b.v b.d;
  (x)

```



# Linear Regression problem

Given some values of an output  $y$  and some values of a dependent parameter  $x$ , assuming that they have a linear dependence find the best fitting line through the observed points



It is mathematically phrased as:

$$Y_n = \text{normal}(c + mX_n, \sigma)$$

# Linear Regression example

```
let lin obs_points ax ay () =  
  let* m = normal 2. 3. in  
  let* c = normal 0. 10. in  
  let* s = exp 5. in  
  for i = 0 to (obs_points-1) do  
    observe (mk ay.(i) -. m*.mk ax.(i) -. c)  
    (logpdf Primitive.(normal 0. (get s)))  
  done ;
```



# Linear Regression example

```
let lin obs_points ax ay () =
```

```
  let* m = normal 2. 3. in
```

```
  let* c = normal 0. 10. in
```

```
  let* s = exp 5. in
```

```
  for i = 0 to (obs_points-1) do
```

```
    observe (mk ay.(i) -. m*.mk ax.(i) -. c)
```

```
    (logpdf Primitive.(normal 0. (get s)))
```

```
  done ;
```

Some priors that  
we may know

# Linear Regression example

```
let lin obs_points ax ay () =
```

```
  let* m = normal 2. 3. in
```

```
  let* c = normal 0. 10. in
```

```
  let* s = exp 5. in
```

Some priors that  
we may know

```
  for i = 0 to (obs_points-1) do
```

```
    observe (mk ay.(i) -. m*.mk ax.(i) -. c)
```

```
    (logpdf Primitive.(normal 0. (get s)))
```

```
  done ;
```

Iterating over all  
observed  $(x_i, y_i)$

# Linear Regression example

```
let lin obs_points ax ay () =
```

```
  let* m = normal 2. 3. in  
  let* c = normal 0. 10. in  
  let* s = exp 5. in
```

Some priors that  
we may know

```
  for i = 0 to (obs_points-1) do
```

```
    observe (mk ay.(i) -. m*.mk ax.(i) -. c)  
    (logpdf Primitive.(normal 0. (get s)))
```


Iterating over all  
observed  $(x_i, y_i)$

```
  done ;
```


Observes that  
 $y_i - mx_i - c$  is  
coming from a normal  
distribution with  
zero-mean




## Linear Regression result




# Linear Regression example



# Linear Regression Result





# Linear Regression Result



## Comparison with Stan


For a linear regression model we compare how fast the means of the slopes converged in Stan versus EffPPL for the same data given to both models

As can be seen Stan converges faster, as it achieved values very close to the mean in  $10^3$  iterations and  $10^{3.5}$  iterations for the linear regression slopes and constants. While EffPPL took ten times more epochs to achieve a similar proximity to the means of slopes and constants.



## Comparison with Stan


We now compare how close EffPPL was in comparison to Stan in terms of time taken by both libraries to perform a fixed number of epochs. We see that EffPPL performs well as it was able to be roughly 1.5-2X slower than Stan.



# Dimensionality Evaluation

We used a linked list as the main data structure in our implementation. This leads to a complexity of  $O(d^2)$ , where  $d$  are the number of dimensions/parameters used in the model. Use of more advanced data structures can lead to a time complexity of  $O(d)$ .

Plot depicts EffPPL's performance against for varying dimensions. We don't compare against Stan here, because Stan was performing much better at  $O(d)$ .



# Thank you.

Arnhav Datar

[cs18b003@smail.iitm.ac.in](mailto:cs18b003@smail.iitm.ac.in)

+91 9665561490

